

Auto-Encoding Variational Bayes

서상우, 유재창

Generative Model

- ▶ 이 논문에서 소개할 Variational Auto-Encoder는 일종의 Generative Model
- ▶ Generative Model은 입력 변수(*latent variable*) z 로부터 결과물 x (가장 흔하게는 image)을 만들어내는 모델

Autoencoder?

Statisticians additionally have $P_{\theta}(X|Z)$





Plate notation
of a graphical
model
to show off



Maximum likelihood

$$\arg \max_{\theta} \left[p_{\theta}(x) = \int_z p_{\theta}(x, z) = \int_z p_{\theta}(x|z)p_{\theta}(z) \right]$$

$$\int_z q(z|x) = 1$$


$$p(x) = \frac{p(z, x)}{p(z|x)}$$

$$\mathbb{E}[X] = \int x f(x)$$

$$D_{KL}(P||Q) = \int p(x) \log \frac{p(x)}{q(x)}$$

Variational Inference

- ▶ Variational inference의 기본 아이디어는 우리가 posterior inference를 어떻게 할 지 알고 있는 모델 q_ϕ 를 가지고 inference를 하되 parameter ϕ 를 잘 조정해서 P에 최대한 가깝게 만들자.



Deterministic approximation procedures
with *bounds* on probabilities of interest.

Fit the *variational parameters*.

$$q_\phi(z|x) \approx p_\theta(z|x)$$

Kullback Leibler divergence

- ▶ 두 분포가 "가깝다"는 것을 어떻게 표현??
- ▶ 두 분포의 차이를 계산해주는 **KL divergence!**
- ▶ Z 가 연속하면 적분으로 표현가능

$$KL(Q_{\phi}(Z|X)||P(Z|X)) = \sum_{z \in Z} q_{\phi}(z|x) \log \frac{q_{\phi}(z|x)}{p(z|x)}.$$

Variational Inference

$$\log(p(x)) = \int \log(p(x)) q_{\phi}(z|x) dz \quad \leftarrow \int q_{\phi}(z|x) dz = 1$$

$$= \int \log\left(\frac{p(x, z)}{p(z|x)}\right) q_{\phi}(z|x) dz \quad \leftarrow p(x) = \frac{p(x, z)}{p(z|x)}$$


$$= \int \log\left(\frac{p(x, z)}{q_{\phi}(z|x)} \cdot \frac{q_{\phi}(z|x)}{p(z|x)}\right) q_{\phi}(z|x) dz$$

$$= \underbrace{\int \log\left(\frac{p(x, z)}{q_{\phi}(z|x)}\right) q_{\phi}(z|x) dz}_{ELBO(\phi)} + \underbrace{\int \log\left(\frac{q_{\phi}(z|x)}{p(z|x)}\right) q_{\phi}(z|x) dz}_{KL(q_{\phi}(z|x) \parallel p(z|x))}$$

$ELBO(\phi)$

$KL(q_{\phi}(z|x) \parallel p(z|x))$

두 확률분포 간의 거리 ≥ 0



KL을 최소화하는 $q_{\phi}(z|x)$ 의 ϕ 값을 찾으려 하는데 $p(z|x)$ 를 모르기 때문에, KL최소화 대신에 ELBO를 최대화하는 ϕ 값을 찾는다.

Variational Inference - ELBO

$$\log(p(x)) = ELBO(\phi) + KL(q_\phi(z|x)|p(z|x))$$

$$q_{\phi^*}(z|x) = \operatorname{argmax}_{\phi} ELBO(\phi)$$

$$ELBO(\phi) = \int \log\left(\frac{p(x, z)}{q_\phi(z|x)}\right) q_\phi(z|x) dz$$

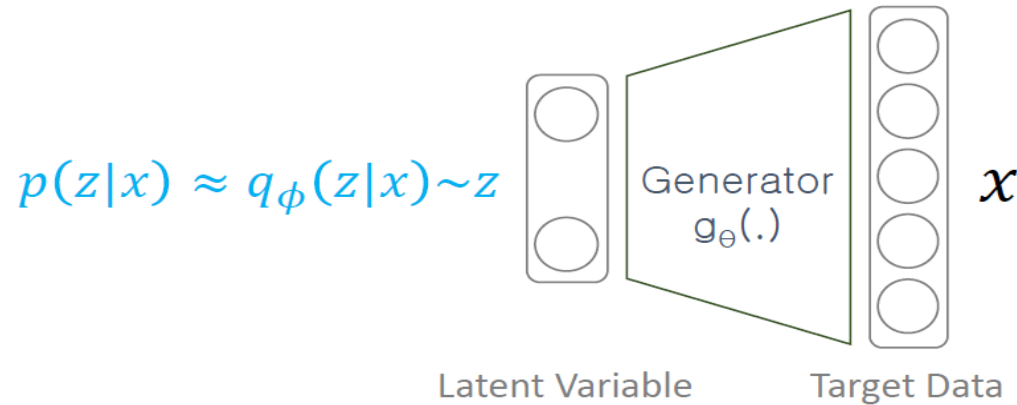
$$= \int \log\left(\frac{p(x|z)p(z)}{q_\phi(z|x)}\right) q_\phi(z|x) dz$$

$$= \int \log(p(x|z)) q_\phi(z|x) dz - \int \log\left(\frac{q_\phi(z|x)}{p(z)}\right) q_\phi(z|x) dz$$

$$= \mathbb{E}_{q_\phi(z|x)}[\log(p(x|z))] - KL(q_\phi(z|x)||p(z))$$
 앞 슬라이드에서의 KL과 인자가 다른 것에 유의

식을 위와 같이 변형하면
KL부분을 쉽게 구할 수 있다.

Variational Inference - loss



Optimization Problem 1 on ϕ : Variational Inference

$$\log(p(x)) \geq \mathbb{E}_{q_\phi(z|x)}[\log(p(x|z))] - KL(q_\phi(z|x)||p(z)) = ELBO(\phi)$$

Optimization Problem 2 on θ : Maximum likelihood

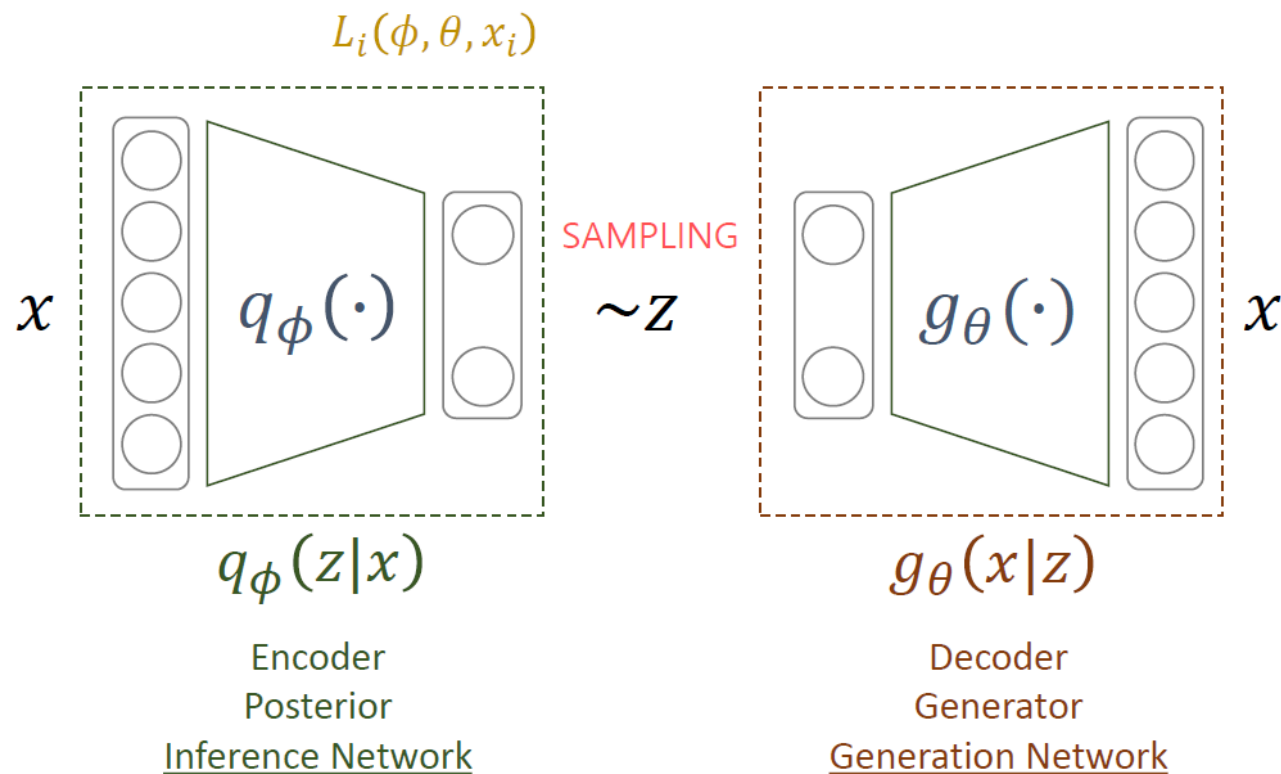
$$-\sum_i \log(p(x_i)) \leq -\sum_i \left\{ \mathbb{E}_{q_\phi(z|x_i)}[\log(p(x_i|g_\theta(z)))] - KL(q_\phi(z|x_i)||p(z)) \right\}$$

Final Optimization Problem

$$\arg \min_{\phi, \theta} \sum_i -\mathbb{E}_{q_\phi(z|x_i)}[\log(p(x_i|g_\theta(z)))] + KL(q_\phi(z|x_i)||p(z))$$

Variational Inference

$$\arg \min_{\phi, \theta} \sum_i \underbrace{-\mathbb{E}_{q_{\phi}(z|x_i)} [\log(p(x_i|g_{\theta}(z)))] + KL(q_{\phi}(z|x_i)||p(z))}_{L_i(\phi, \theta, x_i)}$$



The mathematical basis of VAEs actually has relatively little to do with classical autoencoders

Variational Inference

$$\arg \min_{\phi, \theta} \sum_i \underbrace{-\mathbb{E}_{q_{\phi}(z|x_i)}[\log(p(x_i|g_{\theta}(z)))] + KL(q_{\phi}(z|x_i)||p(z))}_{L_i(\phi, \theta, x_i)}$$

원 데이터에 대한 likelihood

Variational inference를 위한
approximation class 중 선택

다루기 쉬운 확률 분포 중 선택

$$L_i(\phi, \theta, x_i) = \underbrace{-\mathbb{E}_{q_{\phi}(z|x_i)}[\log(p(x_i|g_{\theta}(z)))]}_{\text{Reconstruction Error}} + \underbrace{KL(q_{\phi}(z|x_i)||p(z))}_{\text{Regularization}}$$

Reconstruction Error

- 현재 샘플링 용 함수에 대한 negative log likelihood
- x_i 에 대한 복원 오차 (AutoEncoder 관점)

Regularization

- 현재 샘플링 용 함수에 대한 추가 조건
- 샘플링의 용의성/생성 데이터에 대한 통제성을 위한 조건을 prior에 부여 하고 이와 유사해야 한다는 조건을 부여

Regularization 부분계산

- ▶ Assumption 1


- ▶ $q_\phi \triangleq$ Gaussian distrinution

$$q_\phi(z|x_i) \sim N(\mu_i, \sigma_i^2 I)$$

- ▶ Assumption 2

- ▶ $P(z) \triangleq$ normal distribution

$$p(z) \sim N(0, I)$$



Regularization 부분계산

$$\begin{aligned} KL(q_\phi(z|x_i)||p(z)) &= \frac{1}{2} \left\{ \text{tr}(\sigma_i^2 I) + \mu_i^T \mu_i - J + \ln \frac{1}{\prod_{j=1}^J \sigma_{i,j}^2} \right\} \\ &= \frac{1}{2} \left\{ \sum_{j=1}^J \sigma_{i,j}^2 + \sum_{j=1}^J \mu_{i,j}^2 - J - \sum_{j=1}^J \ln(\sigma_{i,j}^2) \right\} \\ &= \frac{1}{2} \sum_{j=1}^J (\mu_{i,j}^2 + \sigma_{i,j}^2 - \ln(\sigma_{i,j}^2) - 1) \text{ Easy to compute!!} \end{aligned}$$

Kullback–Leibler divergence [\[edit\]](#)

The Kullback–Leibler divergence from $\mathcal{N}_0(\boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0)$ to $\mathcal{N}_1(\boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1)$, for non-singular matrices $\boldsymbol{\Sigma}_0$ and $\boldsymbol{\Sigma}_1$, is:^[8]

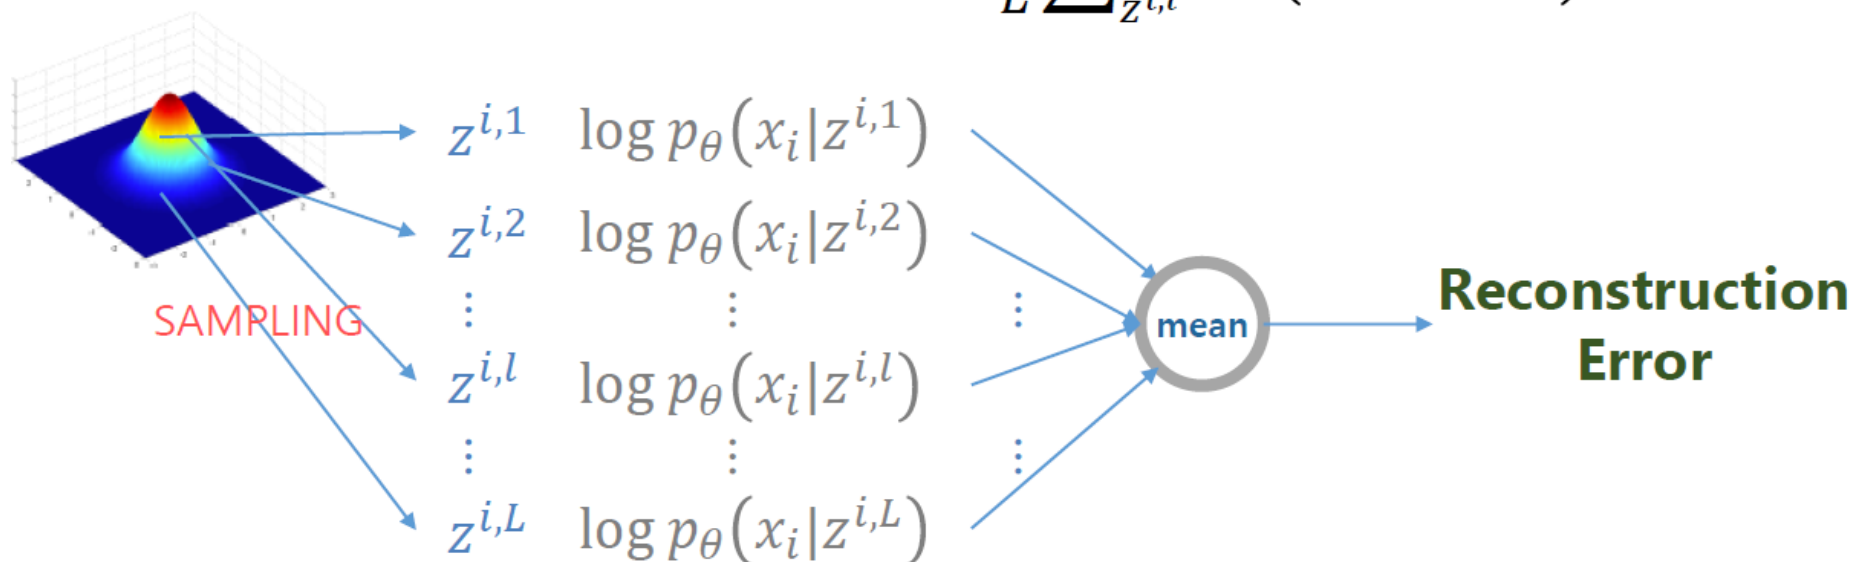
$$D_{\text{KL}}(\mathcal{N}_0||\mathcal{N}_1) = \frac{1}{2} \left\{ \text{tr}(\overset{\text{posterior}}{\boldsymbol{\Sigma}_1^{-1} \boldsymbol{\Sigma}_0}) + (\boldsymbol{\mu}_1 - \overset{\text{prior}}{\boldsymbol{\mu}_0})^T \boldsymbol{\Sigma}_1^{-1} (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_0) - k + \ln \frac{|\boldsymbol{\Sigma}_1|}{|\boldsymbol{\Sigma}_0|} \right\},$$

where k is the dimension of the vector space.

Reconstruction 부분계산

$$\mathbb{E}_{q_{\phi}(z|x_i)}[\log(p_{\theta}(x_i|z))] = \int \log(p_{\theta}(x_i|z))q_{\phi}(z|x_i)dz$$

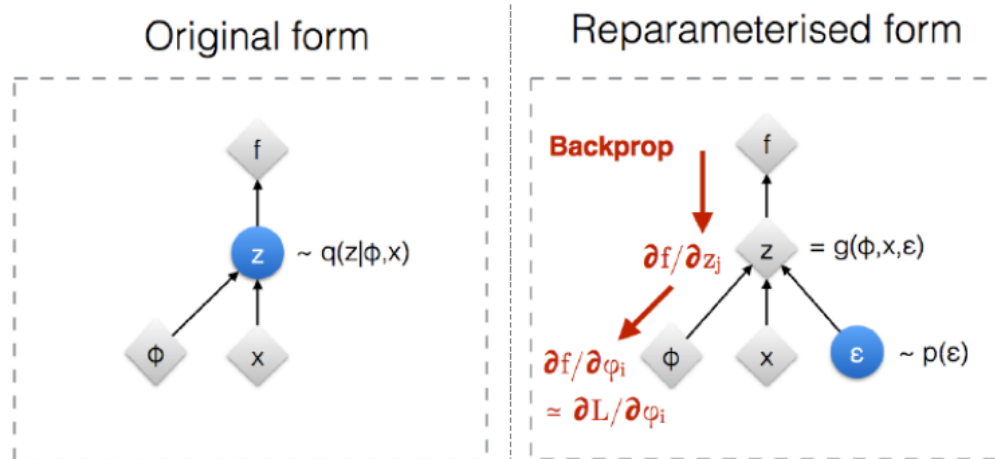
Monte-carlo technique $\rightarrow \approx \frac{1}{L} \sum_{z^{i,l}} \log(p_{\theta}(x_i|z^{i,l}))$





- L is the number of samples for latent vector
- Usually L is set to 1 for convenience

Reconstruction 부분계산

- ▶ 근데 Backpropagation이 불가능...




 : Deterministic node
 : Random node

[Kingma, 2013]
 [Bengio, 2013]
 [Kingma and Welling 2014]
 [Rezende et al 2014]

$$z^{i,l} \sim N(\mu_i, \sigma_i^2 I) \quad \Rightarrow \quad z^{i,l} = \mu_i + \sigma_i^2 \odot \epsilon$$

$$\epsilon \sim N(0, I)$$



Reconstruction 부분계산

$$\int \log(p_{\theta}(x_i|z))q_{\phi}(z|x_i)dz \approx \frac{1}{L} \sum_{z^{i,l}} \log(p_{\theta}(x_i|z^{i,l})) \approx \log(p_{\theta}(x_i|z^i))$$

Monte-carlo technique L=1

Assumption 3-2

[Decoder, likelihood]

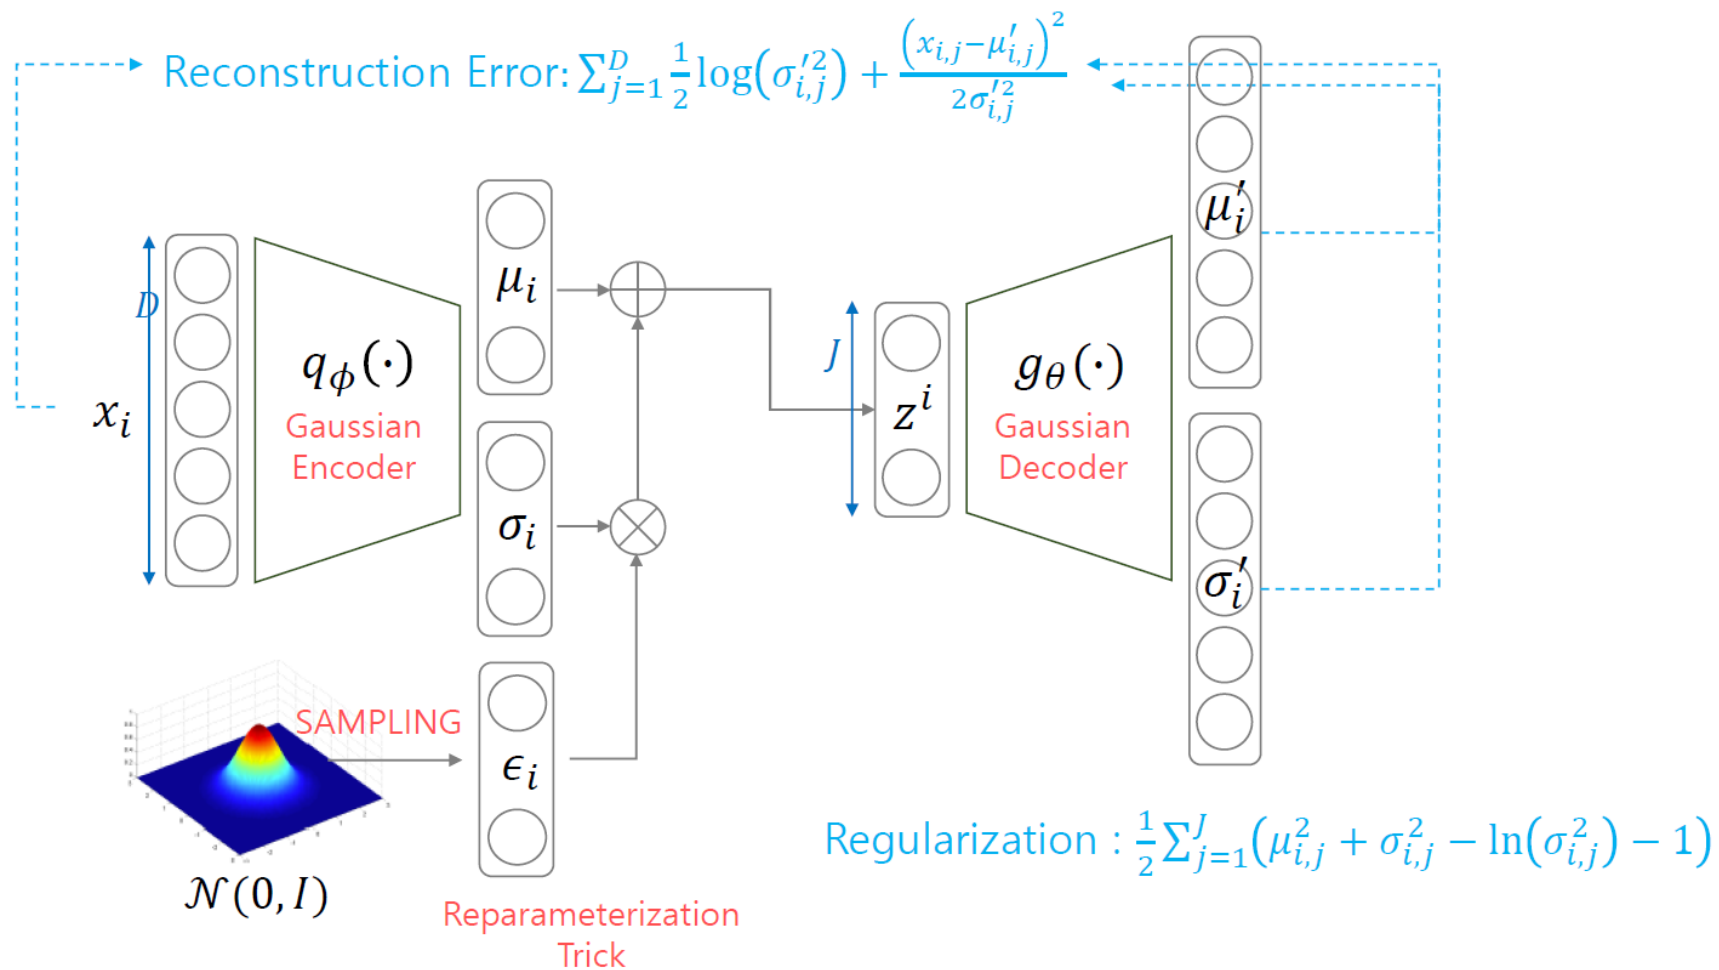
multivariate bernoulli_or gaussian distribution

$$\begin{aligned} \log(p_{\theta}(x_i|z^i)) &= \log(N(x_i; \mu_i, \sigma_i^2 I)) \\ &= - \sum_{j=1}^D \frac{1}{2} \log(\sigma_{i,j}^2) + \frac{(x_{i,j} - \mu_{i,j})^2}{2\sigma_{i,j}^2} \end{aligned}$$

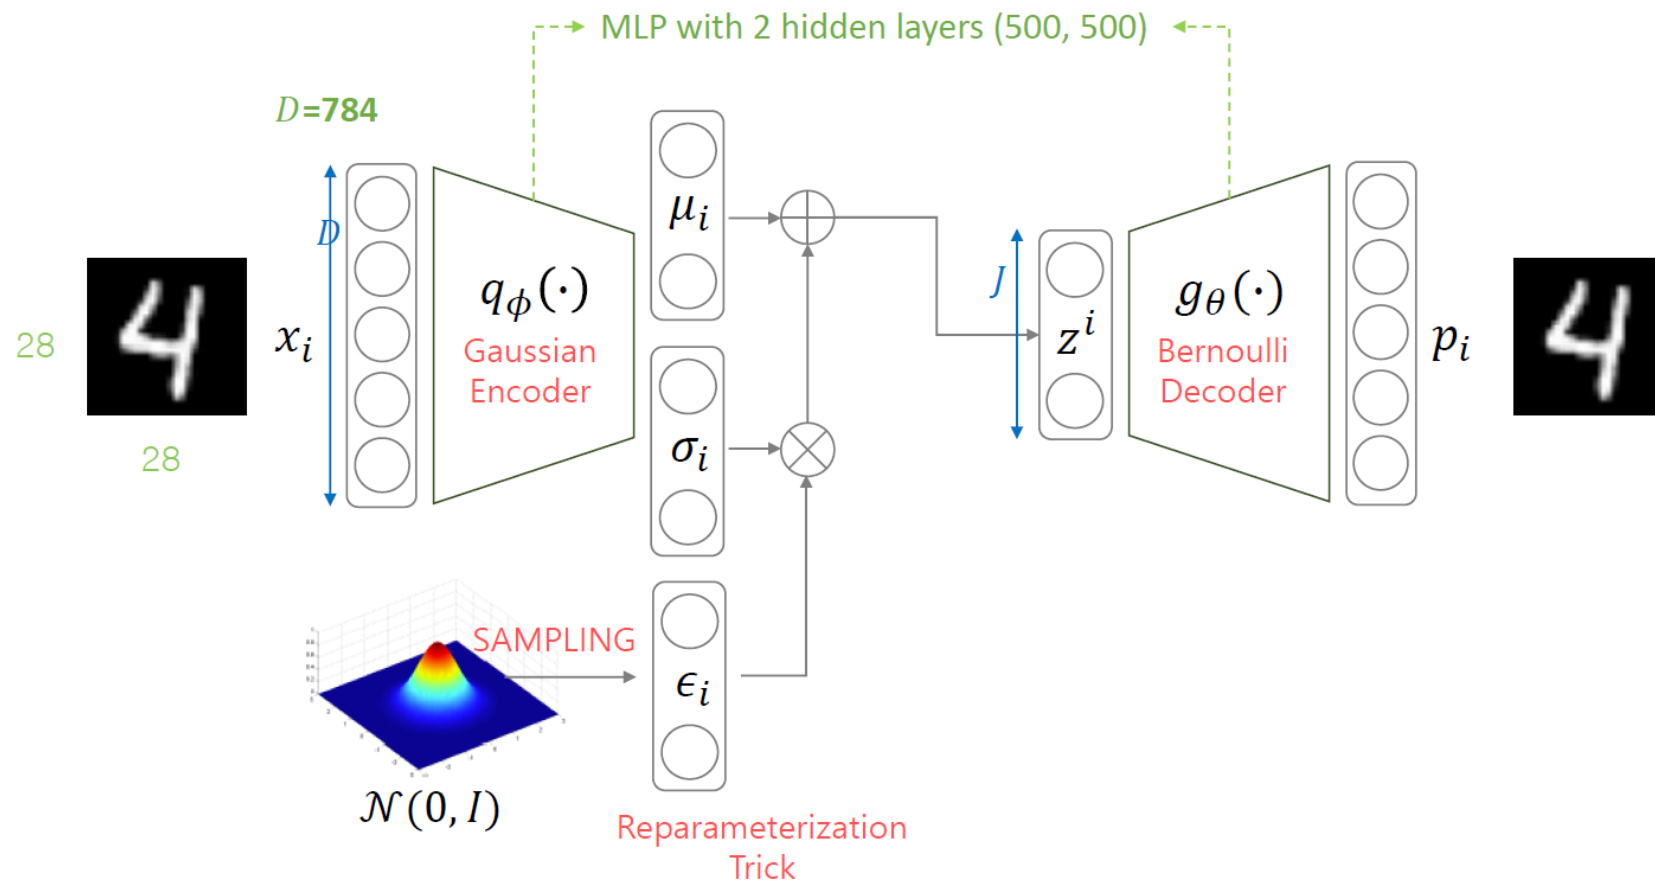
For gaussian distribution with identity covariance

$$\log(p_{\theta}(x_i|z^i)) \propto - \sum_{j=1}^D (x_{i,j} - \mu_{i,j})^2 \quad \leftarrow \text{Squared Error}$$





전체 구조



전체 구조




실험 결과




실험 결과

AE



VAE



test image enc-dec

random latent vector

latent space walking